Final Project – Paper List

For each student, please select one paper from the following list, and make a presentation about that paper. The slides should be written in English.
Notice: Each paper can only be selected by one student. Please contact Jinmei Liu (email: mg20150005@smail.nju.edu.cn) for selecting your paper. First come first serve.

1. Grandmaster Level in Star- Craft II using Multi-Agent Reinforcement Learning, *Nature* 2019
2. AlphaStar: Mastering the Real-Time Strategy Game StarCraft II, *Nature* 2019.
3. Learning Agile and Dynamic Motor Skills for Legged Robots, *Science* 2019
4. SUNRISE: A Simple Unified Framework for Ensemble Learning in Deep Reinforcement Learning, *ICML* 2021
5. Reinforcement Learning with Prototypical Representations, *ICML* 2021
6. Deep Reinforcement Learning amidst Continual Structured Non-Stationarity, *ICML* 2021
7. UneVEn: Universal Value Exploration for Multi-Agent Reinforcement Learning, *ICML* 2021
8. Multi-Task Reinforcement Learning with Context-based Representations, *ICML* 2021
9. MetaCURE: Meta Reinforcement Learning with Empowerment-Driven Exploration, *ICML* 2021
10. Provably Efficient Reinforcement Learning for Discounted MDPs with Feature Mapping, *ICML* 2021
11. The Emergence of Individuality in Multi-Agent Reinforcement Learning, *ICML* 2021
12. Offline Meta-Reinforcement Learning with Advantage Weighting, *ICML* 2021
13. Tesseract: Tensorised Actors for Multi-Agent Reinforcement Learning, *ICML* 2021
14. Scaling Multi-Agent Reinforcement Learning with Selective Parameter Sharing, *ICML* 2021
15. Offline Reinforcement Learning with Pseudometric Learning, *ICML* 2021
16. Exploration in Approximate Hyper-State Space for Meta Reinforcement Learning, *ICML* 2021
17. Multi-Agent Training beyond Zero-Sum with Correlated Equilibrium Meta-

Solvers, *ICML* 2021

18. Counterfactual Credit Assignment in Model-Free Reinforcement Learning, *ICML* 2021
19. Learning Fair Policies in Decentralized Cooperative Multi-Agent Reinforcement Learning, *ICML* 2021
20. Hierarchical Reinforcement Learning for Integrated Recommendation, *AAAI* 2021
21. Lipschitz Lifelong Reinforcement Learning, *AAAI* 2021
22. Learning Task-–Distribution Reward Shaping with Meta-–Learning, *AAAI* 2021
23. DEAR: Deep Reinforcement Learning for Online Advertising Impression in Recommender Systems, *AAAI* 2021
24. Bayesian Optimized Monte Carlo Planning, *AAAI* 2021
25. Towards Effective Context for Meta-Reinforcement Learning: an Approach based on Contrastive Learning, *AAAI* 2021
26. Reinforcement Learning with Trajectory Feedback, *AAAI* 2021
27. GLIB: Efficient Exploration for Relational Model-–Based Reinforcement Learning via Goal-Literal Babbling, *AAAI* 2021
28. Improving Sample Efficiency in Model-–Free Reinforcement Learning from Images, *AAAI* 2021
29. Constrained Risk-Averse Markov Decision Processes, *AAAI* 2021
30. The Value-–Improvement Path: Towards Better Representations for Reinforcement Learning, *AAAI* 2021
31. The Value-Improvement Path Towards Better Representations for Reinforcement Learning, *AAAI* 2021
32. Sample Efficient Reinforcement Learning with REINFORCE, *AAAI* 2021
33. A General Offline Reinforcement Learning Framework for Interactive Recommendation, *AAAI* 2021
34. Combining Reinforcement Learning and Constraint Programming for Combinatorial Optimization, *AAAI* 2021
35. Visual Transfer for Reinforcement Learning via Wasserstein Domain Confusion, *AAAI* 2021
36. Evolving Reinforcement Learning Algorithms, *ICLR* 2021
37. Regularized Inverse Reinforcement Learning, *ICLR* 2021
38. Correcting experience replay for multi-agent communication, *ICLR* 2021
39. Transient Non-stationarity and Generalisation in Deep Reinforcement Learning, *ICLR* 2021
40. Data-Efficient Reinforcement Learning with Self-Predictive Representations, *ICLR* 2021
41. Sample-Efficient Automated Deep Reinforcement Learning, *ICLR* 2021

42. Hierarchical Reinforcement Learning by Discovering Intrinsic Options, *ICLR* 2021
43. A Policy Gradient Algorithm for Learning to Learn in Multiagent Reinforcement Learning, *ICLR* 2021
44. Meta-Reinforcement Learning Robust to Distributional Shift via Model Identification and Experience Relabeling, *ICLR* 2021
45. Transfer among Agents: An Efficient Multiagent Transfer Learning Framework, *ICLR* 2021
46. Safe Reinforcement Learning with Natural Language Constraints, *ICLR* 2021
47. Softmax Deep Double Deterministic Policy Gradients, *NIPS* 2020
48. Breaking the Sample Size Barrier in Model-Based Reinforcement Learning with a Generative Model, *NIPS* 2020
49. Learning Individually Inferred Communication for Multi-Agent Cooperation, *NIPS* 2020
50. Almost Optimal Model-Free Reinforcement Learningvia Reference-Advantage Decomposition, *NIPS* 2020
51. Efficient Model-Based Reinforcement Learning through Optimistic Policy Search and Planning, *NIPS* 2020
52. An Improved Analysis of (Variance-Reduced) Policy Gradient and Natural Policy Gradient Methods, *NIPS* 2020
53. MOReL: Model-Based Offline Reinforcement Learning, *NIPS* 2020
54. Safe Reinforcement Learning via Curriculum Induction, *NIPS* 2020
55. Munchausen Reinforcement Learning, *NIPS* 2020
56. A Self-Tuning Actor-Critic Algorithm, *NIPS* 2020
57. Variational Policy Gradient Method for Reinforcement Learning with General Utilities, *NIPS* 2020
58. Reinforcement Learning with Augmented Data, *NIPS* 2020
59. MOPO: Model-based Offline Policy Optimization, *NIPS* 2020
60. Stochastic Latent Actor-Critic: Deep Reinforcement Learning with a Latent Variable Model, *NIPS* 2020