

# Lecture 9: Deep Q-learning

Zhi Wang & Chunlin Chen

Department of Control and Systems Engineering  
Nanjing University

Nov. 26th, 2020

# Contents and Goals

- How we can make Q-learning work with deep networks
  - Use replay buffers, separate target networks
- Tricks for improving Q-learning in practice
  - Double Q-learning, multi-step Q-learning
- Continuous Q-learning methods
- Goals
  - Understand how to implement Q-learning so that it can be used with complex function approximators
  - Understand how to extend Q-learning to continuous actions


# Table of Contents

- 1 Q-learning with deep neural networks
  - Problem 1: Correlated samples – Solution 1: Replay buffers
  - Problem 2: Moving target – Solution 2: Target networks
- 2 Deep deterministic policy gradient (DDPG)
  - Continuous action space
  - Approximate the optimal policy using another network
- 3 Extensions
  - Double Q-learning
  - Multi-step returns
  - Practical tips and examples

# Review: Fitted Q-iteration (FQI)

- Full fitted Q-iteration algorithm. Loop:
  1. collect dataset  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy  $\pi$  loop for  $K$  iterations:
    2. set  $y_i \leftarrow r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
    3. set  $\phi \leftarrow \arg \min_\phi \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$

- Online fitted Q-iteration algorithm. Loop:
  1. observe one sample  $(s_i, a_i, r_i, s'_i)$  using behavior policy  $\pi$
  2. set  $y_i \leftarrow r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
  3. set  $\phi \leftarrow \phi - \alpha \frac{dQ_\phi(s_i, a_i)}{d\phi} (Q_\phi(s_i, a_i) - y_i)$



# Problem 1: Correlated samples in MDPs

- Online fitted Q-iteration algorithm. Loop:

1. take some action  $a_i$  observe  $(s_i, a_i, r_i, s'_i)$
2. set  $y_i \leftarrow r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
3. set  $\phi \leftarrow \phi - \alpha \frac{dQ_\phi(s_i, a_i)}{d\phi} (Q_\phi(s_i, a_i) - y_i)$

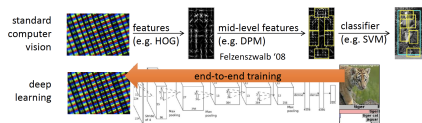
- these samples are correlated!
- Fitted Q-iteration is not gradient descent!

$$\phi \leftarrow \phi - \alpha \frac{dQ_\phi(s_i, a_i)}{d\phi} \left( Q_\phi(s_i, a_i) - \underbrace{\left( r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i) \right)}_{\text{no gradient through target value!}} \right)$$

# Review: Supervised learning vs. Sequential decision-making

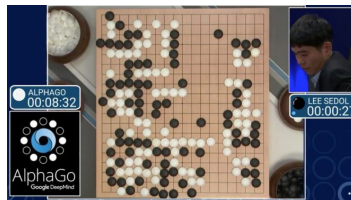
## Supervised learning

- Samples are independent and identically distributed (i.i.d.)
- Given an input, map an optimal output



## Reinforcement learning


- Samples are not i.i.d., temporally correlated
- Given an initial state, find a sequence of optimal actions



# Correlated samples in online Q-learning

- Online fitted Q-iteration algorithm. Loop:
  1. take some action  $a_i$ , observe  $(s_i, a_i, r_i, s'_i)$
  2. set  $\phi \leftarrow \phi - \alpha \frac{dQ_\phi(s_i, a_i)}{d\phi} (Q_\phi(s_i, a_i) - (r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)))$

- sequential states are strongly correlated
- **target value** is always changing



# Correlate samples

- Full fitted Q-iteration algorithm. Loop:
  1. collect **dataset**  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy  $\pi$   
loop for  $K$  iterations:
    2. set  $y_i \leftarrow r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
    3. set  $\phi \leftarrow \arg \min_\phi \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$


- Online fitted Q-iteration. Loop:
  1. take some action  $a_i$ , observe  $(s_i, a_i, r_i, s'_i)$
  2. set  $y_i = r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
  3. set  $\phi \leftarrow \phi - \alpha \frac{dQ_\phi(s_i, a_i)}{d\phi} (Q_\phi(s_i, a_i) - y_i)$

special case with  $K = 1$ ,  
and one gradient step




# How to reduce the correlation between samples?

- Samples in a single episode:
  - temporally correlated




- Samples from different episodes:
  - i.i.d



# Replay buffers: store the data/transitions

- Full fitted Q-iteration algorithm. Loop:
  1. collect **dataset**  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy  $\pi$   
loop for  $K$  iterations:
    2. set  $y_i \leftarrow r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
    3. set  $\phi \leftarrow \arg \min_\phi \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$

- any behavior policy  $\pi$  will work!
- just load data from a **buffer** here
- still use  $K = 1$  and one gradient step




# Q-learning with a replay buffer

- Loop:

1. sample a batch  $\{(s_j, a_j, r_j, s'_j)\}$  from buffer  $\mathcal{B}$
2.  $\phi \leftarrow \phi - \alpha \sum_i \frac{dQ_\phi(s_j, a_j)}{d\phi} \left( Q_\phi(s_j, a_j) - \left( r_j + \gamma \max_{a'_j} Q_\phi(s'_j, a'_j) \right) \right)$

- Step 1: samples are no longer correlated if they come from different episodes
- Step 2: use **multiple samples** in the batch for low-variance gradient
- **Question:** Where does the data come from?
  - Need to periodically feed the replay buffer




# Full Q-learning with a replay buffer

- Loop:

1. collect dataset  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy, add it to  $\mathcal{B}$   
loop for  $K$  iterations:
  2. sample a batch  $\{(s_j, a_j, r_j, s'_j)\}$  from buffer  $\mathcal{B}$
  3.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi(s_j, a_j)}{d\phi} \left( Q_\phi(s_j, a_j) - \left( r_j + \gamma \max_{a'_j} Q_\phi(s'_j, a'_j) \right) \right)$

- $K = 1$  is common, though larger  $K$  is more efficient



# Table of Contents

## 1 Q-learning with deep neural networks

- Problem 1: Correlated samples – Solution 1: Replay buffers
- Problem 2: Moving target – Solution 2: Target networks

## 2 Deep deterministic policy gradient (DDPG)

- Continuous action space
- Approximate the optimal policy using another network

## 3 Extensions

- Double Q-learning
- Multi-step returns
- Practical tips and examples

## Problem 2: Moving target in the Bellman equation

- Online fitted Q-iteration algorithm. Loop:

1. take some action  $a_i$  observe  $(s_i, a_i, r_i, s'_i)$
2. set  $y_i \leftarrow r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
3. set  $\phi \leftarrow \phi - \alpha \frac{dQ_\phi(s_i, a_i)}{d\phi} (Q_\phi(s_i, a_i) - y_i)$

- Samples are correlated: solved by a replay buffer
- Fitted Q-iteration is not gradient descent!

- Target value changes when the Q-network  $\phi$  is updated!

$$\phi \leftarrow \phi - \alpha \frac{dQ_\phi(s_i, a_i)}{d\phi} \left( Q_\phi(s_i, a_i) - \underbrace{\left( r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i) \right)}_{\text{no gradient through target value!}} \right)$$

# The moving target

- Full Q-learning with a replay buffer. Loop:

1. collect dataset  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy, add it to  $\mathcal{B}$  loop for  $K$  iterations:

2. sample a batch  $\{(s_j, a_j, r_j, s'_j)\}$  from buffer  $\mathcal{B}$

3. 
$$\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi(s_j, a_j)}{d\phi} \underbrace{\left( Q_\phi(s_j, a_j) - \left( r_j + \gamma \max_{a'_j} Q_\phi(s'_j, a'_j) \right) \right)}_{\text{one gradient step, moving target}}$$

- Full fitted Q-iteration algorithm. Loop:

1. collect dataset  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy  $\pi$  loop for  $K$  iterations:

2. set  $y_i \leftarrow r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$

3. set 
$$\phi \leftarrow \arg \min_{\phi} \sum_i ||Q_\phi(s_i, a_i) - y_i||^2$$

perfectly well-defined, stable regression

## Solution 2: Target networks

- Idea: use another Q-network and fix it in the inner loop
  - Targets don't change in the inner loop

Q-learning with replay buffer and target network. Loop:

1. save target network parameters:  $\phi' \leftarrow \phi$

loop for  $N$  iterations:

2. collect dataset  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy, add it to  $\mathcal{B}$

loop for  $K$  iterations:

3. sample a batch  $\{(s_j, a_j, r_j, s'_j)\}$  from buffer  $\mathcal{B}$

4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi(s_j, a_j)}{d\phi} \left( Q_\phi(s_j, a_j) - \left( r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j) \right) \right)$



# “Classic” deep Q-network (DQN)

Q-learning with replay buffer and target network. Loop:

1. save target network parameters:  $\phi' \leftarrow \phi$

loop for  $N$  iterations:

2. collect dataset  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy, add it to  $\mathcal{B}$

loop for  $K$  iterations:

3. sample a batch  $\{(s_j, a_j, r_j, s'_j)\}$  from buffer  $\mathcal{B}$

4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_{\phi}(s_j, a_j)}{d\phi} \left( Q_{\phi}(s_j, a_j) - \left( r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j) \right) \right)$

• Classic deep Q-learning with  $K = 1$ . Loop:

1. take some action  $a_i$  and observe  $(s_i, a_i, s'_i, r_i)$ , add it to  $\mathcal{B}$

2. sample mini-batch  $\{(s_j, a_j, r_j, s'_j)\}$  from  $\mathcal{B}$  uniformly

3. compute  $y_j = r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j)$  using target network  $Q_{\phi'}$

4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_{\phi}(s_j, a_j)}{d\phi} (Q_{\phi}(s_j, a_j) - y_j)$

5. update  $\phi'$ : copy  $\phi$  every  $N$  steps



# Alternative target network

- Classic deep Q-learning with  $K = 1$ . Loop:
  1. take some action  $a_i$  and observe  $(s_i, a_i, s'_i, r_i)$ , add it to  $\mathcal{B}$
  2. sample mini-batch  $\{(s_j, a_j, r_j, s'_j)\}$  from  $\mathcal{B}$  uniformly
  3. compute  $y_j = r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j)$  using target network  $Q_{\phi'}$
  4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi(s_j, a_j)}{d\phi} (Q_\phi(s_j, a_j) - y_j)$
  5. **update  $\phi'$ : copy  $\phi$  every  $N$  steps**

- Feels weirdly uneven, can we always have the same lag?
- Popular alternative updating for the target network:

5. **update  $\phi'$  :  $\phi' \leftarrow \tau\phi' + (1 - \tau)\phi$**

- $\tau = 0.99$  works well

# Deep Q-learning and fitted Q-iteration

Deep Q-learning ( $N = 1, K = 1$ ). Loop:

1. save target network parameters:  $\phi' \leftarrow \phi$

loop for  $N$  iterations:

2. collect  $M$  transitions  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy, add them to  $\mathcal{B}$

loop for  $K$  iterations:

3. sample a batch  $\{(s_j, a_j, r_j, s'_j)\}$  from buffer  $\mathcal{B}$

4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi(s_j, a_j)}{d\phi} \left( Q_\phi(s_j, a_j) - \left( r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j) \right) \right)$

Fitted Q-iteration (written similarly as above). Loop:

1. collect  $M$  transitions  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy, add them to  $\mathcal{B}$

loop for  $N$  iterations:

2. save target network parameters:  $\phi' \leftarrow \phi$

loop for  $K$  iterations:

3. sample a batch  $\{(s_j, a_j, r_j, s'_j)\}$  from buffer  $\mathcal{B}$

4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi(s_j, a_j)}{d\phi} \left( Q_\phi(s_j, a_j) - \left( r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j) \right) \right)$

# A more general view

Deep Q-learning ( $N = 1, K = 1$ ). Loop:


1. save target network parameters:  $\phi' \leftarrow \phi$

loop for  $N$  iterations:


2. collect  $M$  transitions  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy, add them to  $\mathcal{B}$   
loop for  $K$  iterations:

3. sample a batch  $\{(s_j, a_j, r_j, s'_j)\}$  from buffer  $\mathcal{B}$

4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi(s_j, a_j)}{d\phi} \left( Q_\phi(s_j, a_j) - \left( r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j) \right) \right)$



# A more general view



- Online fitted Q-iteration: evict immediately, process 1, process 2, and process 3 run at the same speed
- DQN: process 1 and process 3 run at the same speed, process 2 is slow
- Fitted Q-iteration: process 3 is in the inner loop of process 2, which is in the inner loop of process 1

# Table of Contents

- 1 Q-learning with deep neural networks
  - Problem 1: Correlated samples – Solution 1: Replay buffers
  - Problem 2: Moving target – Solution 2: Target networks
- 2 Deep deterministic policy gradient (DDPG)
  - Continuous action space
  - Approximate the optimal policy using another network
- 3 Extensions
  - Double Q-learning
  - Multi-step returns
  - Practical tips and examples

# What's the problem with continuous actions?

- Full fitted Q-iteration algorithm. Loop:

1. collect dataset  $\{(s_i, a_i, r_i, s'_i)\}$  using behavior policy  $\pi$   
loop for  $K$  iterations:
  2. set  $y_i \leftarrow r_i + \gamma \max_{a'_i} Q_\phi(s'_i, a'_i)$
  3. set  $\phi \leftarrow \arg \min_\phi \sum_i \|Q_\phi(s_i, a_i) - y_i\|^2$

- Classic deep Q-learning. Loop:

1. take some action  $a_i$  and observe  $(s_i, a_i, s'_i, r_i)$ , add it to  $\mathcal{B}$
2. sample mini-batch  $\{(s_j, a_j, r_j, s'_j)\}$  from  $\mathcal{B}$  uniformly
3. compute  $y_j = r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j)$  using target network  $Q_{\phi'}$
4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi(s_j, a_j)}{d\phi} (Q_\phi(s_j, a_j) - y_j)$
5. update  $\phi'$ : copy  $\phi$  every  $N$  steps



# The target value involves the max operator

- Classic deep Q-learning. Loop:

1. take some action  $a_i$  and observe  $(s_i, a_i, s'_i, r_i)$ , add it to  $\mathcal{B}$
2. sample mini-batch  $\{(s_j, a_j, r_j, s'_j)\}$  from  $\mathcal{B}$  uniformly
3. compute  $y_j = r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j)$  using target network  $Q_{\phi'}$
4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi(s_j, a_j)}{d\phi} (Q_\phi(s_j, a_j) - y_j)$
5. update  $\phi'$ : copy  $\phi$  every  $N$  steps

$$\pi(a|s) = \begin{cases} 1 & \text{if } a = \arg \max_a Q_\phi(s, a) \\ 0 & \text{otherwise} \end{cases}$$

- target value  $y_j = r_j + \gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j)$ 
  - particularly problematic, need another inner loop of optimization
  - **Question:** how to perform the optimization, i.e., the max operator?

# Option 1: Stochastic optimization

- The action space is typically low-dimensional
  - What about stochastic optimization?

The simplest solution: uniform sampling

- $\max_a Q(s, a) \approx \max\{Q(s, a_1), \dots, Q(s, a_n)\}$
- $(a_1, \dots, a_n)$  sampled from the some distribution (e.g., uniform)

+ dead simple  
+ efficiently parallelizable  
-not very accurate

# More accurate solution: Cross-entropy method (CEM)

Simple **iterative** stochastic optimization:

1. Draw a sample from a probability distribution
2. Minimize the cross-entropy between this distribution and a target distribution to produce a better sample in the next iteration

works OK, for up to about 40 dimensions

A simple example of maximizing  $f(\mathbf{x})$ . Loop:

1. Obtain  $N$  samples:  $\mathbf{X} \sim \text{SampleGaussian}(\mu, \sigma^2; N)$
2. Evaluate objective function  $f(\mathbf{X})$  at sampled points
3. Sort  $\mathbf{X}$  by  $f(\mathbf{X})$  in descending order:  $\mathbf{X} \leftarrow \text{sort}(\mathbf{X}, f)$
4. Update sampling distribution by the top  $M$  elites:  
 $\mu \leftarrow \text{mean}(\mathbf{X}(1:M)), \quad \sigma^2 \leftarrow \text{var}(\mathbf{X}(1:M))$

Objective:

$$\begin{aligned} \mathbf{x}^* &= \arg \max_{\mathbf{x}} f(\mathbf{x}) \\ &\Downarrow \\ \mathbf{a}^* &= \arg \max_{\mathbf{a}} Q(s, \mathbf{a}) \end{aligned}$$


# Many stochastic optimization solutions...

- Covariance matrix adaptation evolution strategy (CMA-ES)
  - an evolutionary algorithm for difficult non-linear non-convex black-box optimization problems in continuous domain
- Many more solutions...

## Option 2: Easily maximizable Q-functions

- Use function class that is easy to optimize
  - e.g., the quadratic function

$$Q_{\phi}(s, a) = -\frac{1}{2}(a - \mu_{\phi}(s))^T P_{\phi}(s)(a - \mu_{\phi}(s)) + V_{\phi}(s)$$



- **NAF: Normalized Advantage Functions**

$$\arg \max_a Q_{\phi}(s, a) = \mu_{\phi}(s)$$

$$\max_a Q_{\phi}(s, a) = V_{\phi}(s)$$

- + no change to algorithm
- + just as efficient as Q-learning
- loses representational power

# Table of Contents

- 1 Q-learning with deep neural networks
  - Problem 1: Correlated samples – Solution 1: Replay buffers
  - Problem 2: Moving target – Solution 2: Target networks
- 2 Deep deterministic policy gradient (DDPG)
  - Continuous action space
  - Approximate the optimal policy using another network
- 3 Extensions
  - Double Q-learning
  - Multi-step returns
  - Practical tips and examples

## Option 3: learn an approximate maximizer

- Lillicrap et al., “Continuous control with deep reinforcement learning,” ICLR 2016.
  - Deep deterministic policy gradient (DDPG)
  - Really approximate **deep Q-learning** in the continuous action domain
- $\max_a Q_\phi(s, a) = Q_\phi(s, \arg \max_a Q_\phi(s, a))$
- **idea**: train another network  $\mu_\theta(s)$  such that

$$\mu_\theta(s) \approx \arg \max_a Q_\phi(s, a)$$

- **Question**: how to optimize this deterministic “actor”  $\mu_\theta(s)$ ?

# Q-learning with continuous actions

- **idea**: train another network  $\mu_\theta(s)$  such that

$$\mu_\theta(s) \approx \arg \max_a Q_\phi(s, a)$$

- how? just solve  $\theta \leftarrow \arg \max_\theta Q_\phi(s, \mu_\theta(s))$

$$\frac{dQ_\phi(s, \mu_\theta(s))}{d\theta} = \frac{dQ_\phi}{da} \cdot \frac{da}{d\theta} = \frac{dQ_\phi}{d\mu_\theta(s)} \cdot \frac{d\mu_\theta(s)}{d\theta}$$

- new target

$$y_j = r_j + \gamma Q_{\phi'}(s'_j, \mu_\theta(s'_j)) \approx r_j + \gamma Q_{\phi'}(s'_j, \arg \max_{a'_j} Q_{\phi'}(s'_j, a'_j))$$




# DDPG network architecture

- Backpropagate the critic:

$$\nabla_{\phi} = \frac{dQ_{\phi}(s, a)}{d\phi} (Q_{\phi}(s, a) - y)$$

- Backpropagate the actor:

$$\nabla_{\theta} = \frac{dQ_{\phi}}{d\mu_{\theta}(s)} \cdot \frac{d\mu_{\theta}(s)}{d\theta}$$



# Deep deterministic policy gradient (DDPG)

- Loop:

1. take some action  $a_i$  and observe  $(s_i, a_i, s'_i, r_i)$ , add it to  $\mathcal{B}$
2. sample mini-batch  $\{(s_j, a_j, r_j, s'_j)\}$  from  $\mathcal{B}$  uniformly
3. compute  $y_j = r_j + \gamma Q_{\phi'}(s'_j, \mu_{\theta'}(s'_j))$  by target networks  $Q_{\phi'}$  and  $\mu_{\theta'}$
4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_{\phi}(s_j, a_j)}{d\phi} (Q_{\phi}(s_j, a_j) - y_j)$
5.  $\theta \leftarrow \theta + \beta \sum_j \frac{dQ_{\phi}}{d\mu_{\theta}(s_j)} \frac{d\mu_{\theta}(s_j)}{d\theta}$
6. update  $\phi', \theta'$ :  $\phi' \leftarrow \tau \phi' + (1 - \tau)\phi$ ,  $\theta' \leftarrow \tau \theta' + (1 - \tau)\theta$

- The behavior policy  $\pi$ :

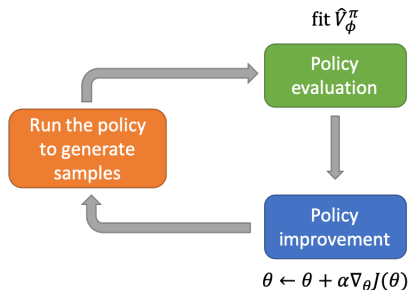
- The target greedy policy is  $\pi^*(s) = \mu_{\theta}(s)$ , actually
- Add some exploration noise to the target greedy policy, just like  $\epsilon$ -greedy in tabular Q-learning

$$\pi(a|s) \sim \mathcal{N}(\mu_{\theta}(s), \sigma^2)$$

# Review: Actor-critic algorithms

- Loop:

1. sample  $\{s_i, a_i, r_i, s'_i\}$  from  $\pi_\theta(a|s)$  (run it on the robot)
2. **policy evaluation**: fit  $\hat{V}_\phi^\pi(s)$  to sampled reward sums
3. evaluate  $\hat{A}^\pi(s_i, a_i) = r_i + \gamma \hat{V}_\phi^\pi(s'_i) - \hat{V}_\phi^\pi(s_i)$
4. **policy improvement**:  $\nabla_\theta J(\theta) \approx \sum_i \nabla_\theta \log \pi_\theta(a_i|s_i) \hat{A}^\pi(s_i, a_i)$
5.  $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$



# DDPG vs. Actor-critic

## • DDPG. Loop:

1. take some action  $a_i$  and observe  $(s_i, a_i, s'_i, r_i)$ , add it to  $\mathcal{B}$
2. sample mini-batch  $\{(s_j, a_j, r_j, s'_j)\}$  from  $\mathcal{B}$  uniformly
3. compute  $y_j = r_j + \gamma Q_{\phi'}(s'_j, \mu_{\theta'}(s'_j))$  by target networks  $Q_{\phi'}$  and  $\mu_{\theta'}$
4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_{\phi}(s_j, a_j)}{d\phi} (Q_{\phi}(s_j, a_j) - y_j)$
5.  $\theta \leftarrow \theta + \beta \sum_j \frac{dQ_{\phi}}{d\mu_{\theta}(s_j)} \frac{d\mu_{\theta}(s_j)}{d\theta}$
6. update  $\phi', \theta'$ :  $\phi' \leftarrow \tau\phi' + (1 - \tau)\phi$ ,  $\theta' \leftarrow \tau\theta' + (1 - \tau)\theta$

## • Actor-critic. Loop:

1. sample  $\{s_i, a_i, r_i, s'_i\}$  from  $\pi_{\theta}(a|s)$  (run it on the robot)
2. policy evaluation: fit  $\hat{V}_{\phi}^{\pi}(s)$  to sampled reward sums
3. evaluate  $\hat{A}^{\pi}(s_i, a_i) = r_i + \gamma \hat{V}_{\phi}^{\pi}(s'_i) - \hat{V}_{\phi}^{\pi}(s_i)$
4. policy improvement:  $\nabla_{\theta} J(\theta) \approx \sum_i \nabla_{\theta} \log \pi_{\theta}(a_i | s_i) \hat{A}^{\pi}(s_i, a_i)$
5.  $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$

# Review: Q-learning vs. SARSA

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

- Q-learning approximates the optimal action-value function for an optimal policy,  $Q \approx Q_* = Q_{\pi_*}$ 
    - The target policy is greedy w.r.t  $Q$ ,  $\pi(a|s) = \arg \max_a Q(s, a)$
    - The behavior policy can be others, e.g.,  $b(a|s) = \varepsilon$ -greedy
- 

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

- SARSA approximates the action-value function for the behavior policy,  $Q \approx Q_\pi = Q_b$ 
  - The target and the behavior policy are the same, e.g.,  $\pi(a|s) = b(a|s) = \varepsilon$ -greedy

# DDPG vs. Actor-critic

- DDPG

- The actor: approximate the optimal policy
$$a^* \mu_\theta(s) = \arg \max_a Q_\phi(s, a)$$
- The critic: approximate the optimal action-value function  $Q_\phi^*$
- Off-policy, more sample efficient

- Actor-critic

- The actor: approximate the current policy  $a \sim \pi_\theta(a|s)$
- The critic: approximate the state-value function  $V_\phi^\pi$  for given policy  $\pi$
- On-policy, at least converge to a local optimum

# Table of Contents

- 1 Q-learning with deep neural networks
  - Problem 1: Correlated samples – Solution 1: Replay buffers
  - Problem 2: Moving target – Solution 2: Target networks
- 2 Deep deterministic policy gradient (DDPG)
  - Continuous action space
  - Approximate the optimal policy using another network
- 3 Extensions
  - Double Q-learning
  - Multi-step returns
  - Practical tips and examples

# Overestimation in Q-learning

- target value  $y_j = r_j + \underbrace{\gamma \max_{a'_j} Q_{\phi'}(s'_j, a'_j)}_{\text{this is the problem}}$

- Imagine we have two random variables:  $x_1$  and  $x_2$

$$\mathbb{E}[\max(x_1, x_2)] \geq \max(\mathbb{E}[x_1], \mathbb{E}[x_2])$$

- $Q_{\phi'}(s', a')$  is not perfect – it looks “noisy”
- hence  $\max_{a'} Q_{\phi'}(s', a')$  **overestimates** the next value!
- note that  $\max_{a'} Q_{\phi'}(s', a') = Q_{\phi'}(s', \arg \max_{a'} Q_{\phi'}(s', a'))$ 
  - action selected according to  $Q_{\phi'}$
  - value also comes from  $Q_{\phi'}$



# Double Q-learning

- $\mathbb{E}[\max(x_1, x_2)] \geq \max(\mathbb{E}[x_1], \mathbb{E}[x_2])$
- note that  $\max_{a'} Q_{\phi'}(s', a') = Q_{\phi'}(s', \arg \max_{a'} Q_{\phi'}(s', a'))$ 
  - action selected according to  $Q_{\phi'}$
  - value also comes from  $Q_{\phi'}$
  - if the noise in the two parts is decorrelated, the problem goes away!
- **IDEA:** don't use the same network to choose the action and evaluate value!
- “double” Q-learning: use two networks

$$Q_{\phi_A}(s, a) \leftarrow r + \gamma Q_{\phi_B}(s', \arg \max_{a'} Q_{\phi_A}(s', a'))$$

$$Q_{\phi_B}(s, a) \leftarrow r + \gamma Q_{\phi_A}(s', \arg \max_{a'} Q_{\phi_B}(s', a'))$$

- if the two Q-networks,  $Q_{\phi_A}$  and  $Q_{\phi_B}$ , are noisy in different ways, there is no problem

# Double Q-learning in practice

- Where to get two Q-functions?
  - just use the current and target networks!
- standard Q-learning:  $y = r + \gamma Q_{\phi'}(s', \arg \max_{a'} Q_{\phi'}(s', a'))$
- double Q-learning:  $y = r + \gamma Q_{\phi'}(s', \arg \max_{a'} Q_{\phi}(s', a'))$ 
  - just use **current network** (not target network) to evaluate action
  - still use target network to evaluate value

# Table of Contents

- 1 Q-learning with deep neural networks
  - Problem 1: Correlated samples – Solution 1: Replay buffers
  - Problem 2: Moving target – Solution 2: Target networks
- 2 Deep deterministic policy gradient (DDPG)
  - Continuous action space
  - Approximate the optimal policy using another network
- 3 Extensions
  - Double Q-learning
  - **Multi-step returns**
  - Practical tips and examples


## $n$ -step bootstrapping: Combine MC and one-step TD

- Neither MC or one-step TD is always the best, we generalize both methods so that one can shift from one to the other smoothly as needed to meet the demands of a particular task
- One-step TD: In many applications, one wants to be able to update the action very fast to take into account anything that has changed
- However, bootstrapping works best if it is over a length of time in which a significant and recognizable state change has occurred

$n = 1$	$n$ -step TD	$n = \infty$
TD(0)	$\longleftrightarrow$	MC

# $n$ -step TD prediction

- Perform an update based on an intermediate number of rewards, more than one, but less than all of them until termination



# Recall MC and TD(0) updates

- In MC updates, the target is the **complete return**

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-t+1} R_T$$

$$\begin{aligned} V(S_t) &\leftarrow V(S_t) + \alpha[G_t - V(S_t)] \\ &= V(S_t) + \alpha[R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-t+1} R_T - V(S_t)] \end{aligned}$$

- In TD(0) updates, the target is the **one-step return**

$$G_{t:t+1} = R_{t+1} + \gamma V(S_{t+1})$$

$$\begin{aligned} V(S_t) &\leftarrow V(S_t) + \alpha[G_{t:t+1} - V(S_t)] \\ &= V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)] \end{aligned}$$

# $n$ -step TD update rule

- For  $n$ -step TD, set the target as the  $n$ -**step return**

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$$

- All  $n$ -step returns can be considered approximations to the complete return, truncated after  $n$  steps and then corrected for the remaining missing terms by  $V(S_{t+n})$

$$\begin{aligned} V(S_t) &\leftarrow V(S_t) + \alpha[G_{t:t+n} - V(S_t)] \\ &= V(S_t) + \alpha[R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n}) - V(S_t)] \end{aligned}$$

# Deep Q-learning with $n$ -step bootstrapping

- Q-learning target:  $y_{j,t} = r_{j,t} + \gamma \max_{a'_{j,t+1}} Q_{\phi'}(s'_{j,t+1}, a'_{j,t+1})$ 
  - these are the only values that matter if  $Q_{\phi'}$  is bad!
  - these values are important if  $Q_{\phi'}$  is good
- Construct multi-step targets,  $N$ -step return estimator:

$$y_{j,t} = \sum_{t'=t}^{t+N-1} \gamma^{t'-t} r_{j,t'} + \gamma^N \max_{a'_{j,t+N}} Q_{\phi'}(s'_{j,t+N}, a'_{j,t+N})$$



# Table of Contents

- 1 Q-learning with deep neural networks
  - Problem 1: Correlated samples – Solution 1: Replay buffers
  - Problem 2: Moving target – Solution 2: Target networks
- 2 Deep deterministic policy gradient (DDPG)
  - Continuous action space
  - Approximate the optimal policy using another network
- 3 Extensions
  - Double Q-learning
  - Multi-step returns
  - Practical tips and examples

# Simple practical tips for Q-learning

- Q-learning takes some care to stabilize
  - Test on easy, reliable tasks first, make sure your implementation is correct





Figure: From T. Schaul, J. Quan, I. Antonoglou, and D. Silver. “Prioritized experience replay”. *arXiv preprint arXiv:1511.05952* (2015), Figure 7

- Large replay buffers help improve stability
  - Looks more like fitted Q-iteration
- It takes time, be patient - might be no better than random for a while
- Start with high exploration and gradually move to high exploitation

# Advanced tips for Q-learning

- Bellman error gradients can be big; clip gradients or use Huber loss

$$\mathcal{L}(x) = \begin{cases} x^2/2 & \text{if } |x| \leq \delta \\ \delta|x| - \delta^2/2 & \text{otherwise} \end{cases}$$



- Double Q-learning *helps a lot* in practice, simple and no downsides
- $N$ -step returns also help a lot, but have some downsides
- Schedule exploration (high to low) and learning rates (high to low)
  - Adam optimizer can help too
- Run multiple random seeds, it's very inconsistent between runs

# Q-learning with convolutional networks


- Mnih et al., “Human-level control through deep reinforcement learning,” 2013.
- Use replay buffer and target network
- One-step backup, one gradient step
- Can be improved a lot with double Q-learning (and other tricks)



Space Invaders



Breakout



Carnival



Boxing



Pong



Kung-Fu Master




Skiing



River Raid




Enduro



Alien




Montezuma's Revenge




Ms. PacMan

# Q-learning on a real robot

- Gu et al., “Robot manipulation with deep reinforcement learning and ...,” 2017.
- Continuous actions with NAF (quadratic in actions)
- Use replay buffer and target network
- One-step backup, four gradient steps per simulator step for efficiency
- Parallelized across multiple robots



- Q-learning with deep neural networks
  - Replay buffers
  - Target networks
- Generalized fitted Q-iteration
- Deep deterministic policy network
  - Deep Q-learning for continuous action space
  - Another network for approximating optimal policy
  - Off-policy
- Extensions
  - Double Q-learning
  - Multi-step Q-learning



# Learning objectives of this lecture



- You should be able to...
  - Use deep neural networks to approximate Q-functions, be able to implement deep Q-learning with replay buffers and target networks
  - Use deep deterministic policy gradient for continuous actions
  - Know double Q-learning for addressing the overestimation problem
  - Know deep Q-learning with  $n$ -step returns

- Lecture 8 of CS285 at UC Berkeley, **Deep Reinforcement Learning, Decision Making, and Control**
  - <http://rail.eecs.berkeley.edu/deeprlcourse/static/slides/lec-8.pdf>
- DRL Q-learning papers
  - Mnih et al. (2013). **Human level control through deep reinforcement learning**: Q-learning with convolutional networks for playing Atari.
  - Van Hasselt, Guez , Silver. (2015). **Deep reinforcement learning with double Q-learning**: a very effective trick to improve performance of deep Q-learning.
  - Lillicrap et al. (2016). **Continuous control with deep reinforcement learning**: continuous Q-learning with actor network for approximate maximization.
  - Wang, Schaul , Hessel, van Hasselt, Lanctot , de Freitas (2016). **Dueling network architectures for deep reinforcement learning**: separates value and advantage estimation in Q-function.
  - Z. Ren, et al., **Self-Paced Prioritized Curriculum Learning With Coverage Penalty in Deep Reinforcement Learning**, *TNNLS*, 2018.



# Homework 5

- Study the DDPG algorithm in detail
- Implement the DDPG algorithm on problems 1 & 2
  - Problem 1: the point maze navigation, continuous state-action space ( $s, a \in \mathbb{R}^2$ ,  $s \in [-0.5, 0.5]^2$ ,  $a \in [-0.1, 0.1]^2$ )
  - Problem 2: the MuJoCo HalfCheetah, make the robot run forward
  - Compare DDPG with policy gradient and actor-critic algorithms
- Write a report introducing the algorithms and your experimentation
  - Explanations, steps, evaluation results, visualizations...
  - Submit the code and the report to [yuan yang@smail.nju.edu.cn](mailto:yuan yang@smail.nju.edu.cn)



# THE END